## QUERY_HISTORY

| Query Id | Query Text | Database Id | Database Name | Schema Id | Schema Name | Query Type | Session |
|---|---|---|---|---|---|---|---|
| 64ebe479-0501-40aa-0004-ba03064a61ba | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219728 |
| c6a06f47-0501-40aa-0004-ba030649b50a | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219728 |
| 3e43ee1c-0501-405d-0004-ba03064a0bce | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219728 |
| bc8468df-0501-405d-0004-ba03064a0bfe | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219728 |
| 3dde65b0-0400-cf74-0004-ba030550af66 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219660 |
| a03bb09a-0400-d0ee-0004-ba0305519c22 | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219661 |
| 3e60e75d-0400-cfef-0004-ba030550691e | commit | null | null | null | ACME_SCHEMA_BASE | COMMIT | 13304219660 |
| 3a8a94ec-0601-1830-0004-ba03060a4a0a | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219707 |
| f50481fb-0401-65da-0004-ba03069f9326 | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219748 |
| 217ef5af-0501-3ebf-0004-ba030647a606 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219728 |
| 5a6a9fdf-0501-3bf4-0004-ba03063fd362 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219726 |
| 4fd44992-0501-5d60-0004-ba03068dc8f6 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219743 |
| 57f75a83-0400-bb98-0004-ba0305293012 | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219647 |
| 9ef819a2-0400-936e-0004-ba0304dd01b6 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219625 |
| 6b413d1f-0600-f184-0004-ba030585956e | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219681 |
| 61e9f5c8-0600-f1a8-0004-ba030585b382 | show objects like X in schema null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SHOW | 13304219681 |
| 8d62fbe3-0500-6e70-0004-ba03049a6b26 | select * from null.ACME_SCHEMA_BASE | null | null | null | ACME_SCHEMA_BASE | SELECT | 13304219604 |

SUMMARY 5,918,618 rows – 68 columns

## DATABASE_STORAGE_USAGE_HISTORY

| Usage Date | Database Id | Database Name | Deleted | Average Database Bytes | Average Failsafe Bytes |
|---|---|---|---|---|---|
| 2025-03-08 00:00:00 | 53 | FINANCIALS | null | 519966350880 | 3440940 |
| 2025-03-08 00:00:00 | 6 | REPORTING_DB | null | 351497675848 | 31503542310 |
| 2025-03-08 00:00:00 | 41 | FINSERV1 | null | 370134535672 | 30088439223 |
| 2025-03-08 00:00:00 | 57 | PUBLIC | null | 1002609581 | 3440940 |
| 2025-03-08 00:00:00 | 60 | FINSERV2 | null | 85480697168 | 3440940 |
| 2025-03-08 00:00:00 | 5 | RESEARCH_DB | null | 584693082169 | 47521918605 |
| 2025-03-09 00:00:00 | 6 | REPORTING_DB | null | 352986832982 | 31574481126 |
| 2025-03-09 00:00:00 | 41 | FINSERV1 | null | 370115116965 | 30102406378 |
| 2025-03-09 00:00:00 | 57 | PUBLIC | null | 1002580627 | 3438044 |
| 2025-03-09 00:00:00 | 60 | FINSERV2 | null | 85480668215 | 3438044 |
| 2025-03-09 00:00:00 | 53 | FINANCIALS | null | 519966321927 | 3438044 |
| 2025-03-09 00:00:00 | 5 | RESEARCH_DB | null | 584779346869 | 47730846171 |
| 2025-03-10 00:00:00 | 60 | FINSERV2 | null | 85480637382 | 3434961 |
| 2025-03-10 00:00:00 | 57 | PUBLIC | null | 1002549794 | 3434961 |
| 2025-03-10 00:00:00 | 53 | FINANCIALS | null | 519966291094 | 3434961 |
| 2025-03-10 00:00:00 | 6 | REPORTING_DB | null | 374085252432 | 31621434598 |
| 2025-03-10 00:00:00 | 5 | RESEARCH_DB | null | 614723983802 | 47984355274 |

SUMMARY 2,981 rows – 6 columns

## WAREHOUSE_EVENTS_HISTORY

| Timestamp | Warehouse Id | Warehouse Name | Cluster Number | Event Name | Event Reason | Event State | User Name | Role Name | Que |
|---|---|---|---|---|---|---|---|---|---|
| 2024-09-07 08:25:32 | 15 | INVENTORY_SCHEDULING_WH | 0 | RESUME_WAREHOUSE | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:03:04 | 4 | SALES_PROD_WH | 0 | RESUME_CLUSTER | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:28:35 | 4 | SALES_PROD_WH | 0 | SUSPEND_CLUSTER | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-07 08:55:19 | 10 | SMOPS_WH | 0 | RESUME_CLUSTER | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:21:42 | 4 | SALES_PROD_WH | 0 | RESUME_CLUSTER | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:54:16 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-07 08:31:07 | 4 | SALES_PROD_WH | 0 | RESUME_WAREHOUSE | WAREHOUSE_AUTORESUME | STARTED | null | null | null |
| 2024-09-07 08:36:36 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-07 08:46:26 | 4 | SALES_PROD_WH | 0 | RESUME_WAREHOUSE | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:08:26 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-07 08:08:25 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | STARTED | null | null | null |
| 2024-09-07 08:31:08 | 4 | SALES_PROD_WH | 0 | RESUME_CLUSTER | WAREHOUSE_AUTORESUME | COMPLETED | null | null | null |
| 2024-09-07 08:55:18 | 10 | SMOPS_WH | 0 | RESUME_WAREHOUSE | WAREHOUSE_AUTORESUME | STARTED | null | null | null |
| 2024-09-08 07:08:16 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-08 07:08:16 | 4 | SALES_PROD_WH | 0 | SUSPEND_CLUSTER | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |
| 2024-09-08 18:37:12 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | STARTED | null | null | null |
| 2024-09-08 22:08:26 | 4 | SALES_PROD_WH | 0 | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND | COMPLETED | null | null | null |

SUMMARY 125,247 rows – 10 columns

## WAREHOUSE_METERING_HISTORY

| Start Time | End Time | Warehouse Id | Warehouse Name | Credits Used | Credits Used Compute | Credits Used Clou |
|---|---|---|---|---|---|---|
| 2025-01-11 11:00:00 | 2025-01-11 11:00:00 | 10 | SMOPS_WH | 7.322897778 | 7.308130961 | |
| 2025-01-12 05:00:00 | 2025-01-12 06:00:00 | 4 | SALES_PROD_WH | 5.376124345 | 5.3675534 | |
| 2025-01-12 04:00:00 | 2025-01-12 05:00:00 | 4 | SALES_PROD_WH | 4.291854937 | 4.285003132 | |
| 2025-01-11 11:00:00 | 2025-01-11 11:00:00 | 8 | DATA_ENGINEERING | 4.21260231 | 4.205722207 | |
| 2025-01-11 11:00:00 | 2025-01-11 12:00:00 | 10 | SMOPS_WH | 6.882956635 | 6.871896731 | |
| 2025-01-11 11:00:00 | 2025-01-11 12:00:00 | 4 | SALES_PROD_WH | 5.824417617 | 5.807150413 | |
| 2025-01-12 06:00:00 | 2025-01-12 07:00:00 | 4 | SALES_PROD_WH | 7.772005077 | 7.759536124 | |
| 2025-01-11 11:00:00 | 2025-01-11 11:00:00 | 7 | WEBSITE_WH | 6.204036994 | 6.193517946 | |
| 2025-01-11 17:00:00 | 2025-01-11 18:00:00 | 4 | SALES_PROD_WH | 4.72537878 | 4.717844873 | |
| 2025-01-11 14:00:00 | 2025-01-11 15:00:00 | 4 | SALES_PROD_WH | 4.988816742 | 4.980882102 | |
| 2025-01-12 07:00:00 | 2025-01-12 08:00:00 | 4 | SALES_PROD_WH | 5.768946695 | 5.758584458 | |
| 2025-01-11 21:00:00 | 2025-01-11 22:00:00 | 4 | SALES_PROD_WH | 7.035810844 | 7.024720014 | |
| 2025-01-11 09:00:00 | 2025-01-11 11:00:00 | 15 | INVENTORY_SCHEDULING_WH | 6.633668003 | 6.623136803 | |
| 2025-01-12 03:00:00 | 2025-01-12 04:00:00 | 4 | SALES_PROD_WH | 5.019123823 | 5.011119743 | |
| 2025-01-11 12:00:00 | 2025-01-11 13:00:00 | 4 | SALES_PROD_WH | 7.842316254 | 7.827281284 | |
| 2025-01-11 18:00:00 | 2025-01-11 19:00:00 | 4 | SALES_PROD_WH | 5.943806031 | 5.934296529 | |
| 2025-01-11 23:00:00 | 2025-01-12 00:00:00 | 4 | SALES_PROD_WH | 4.578383972 | 4.571233634 | |

SUMMARY 18,178 rows – 7 columns

Data | Query History Analysis | Storage Insights | Events Insights | Credits Insights | Warehouse Efficiency Summary

Analysis Data

# Query Volume and Performance Dashboard

**Start Time**

Last 30 days including today (05/12/2025 – 06/10/2025)

## Total Queries by Warehouse Id

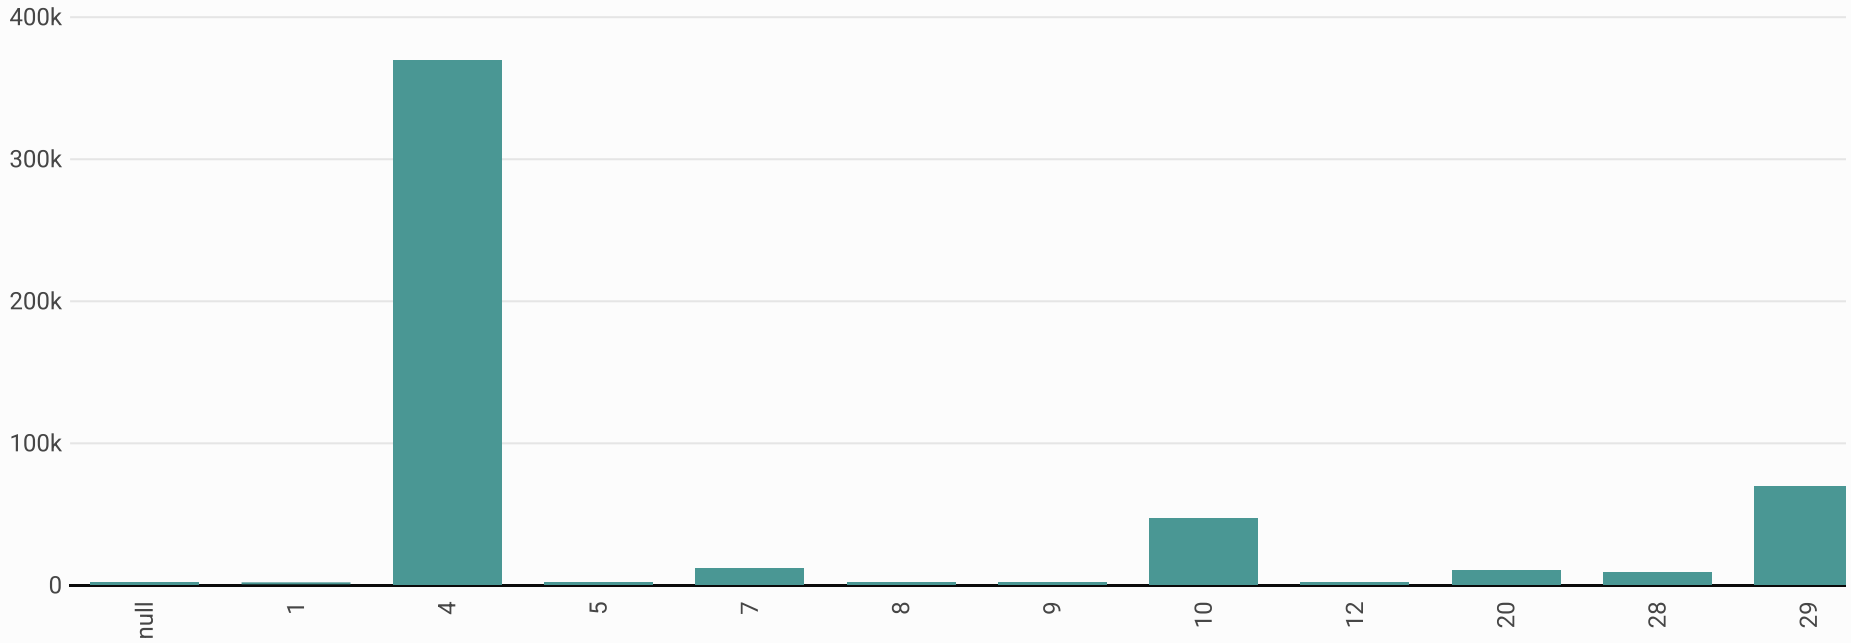


**Key Takeaways:**

**Top Performer:** The tallest bar is the busiest warehouse. This is where most of our compute demand lives.

**Cold spots:** Any warehouses with near-zero bars indicate under-utilized capacity, can be resized.

---

**Start Time**

Last 30 days including today (05/12/2025 – 06/10/2025) ×

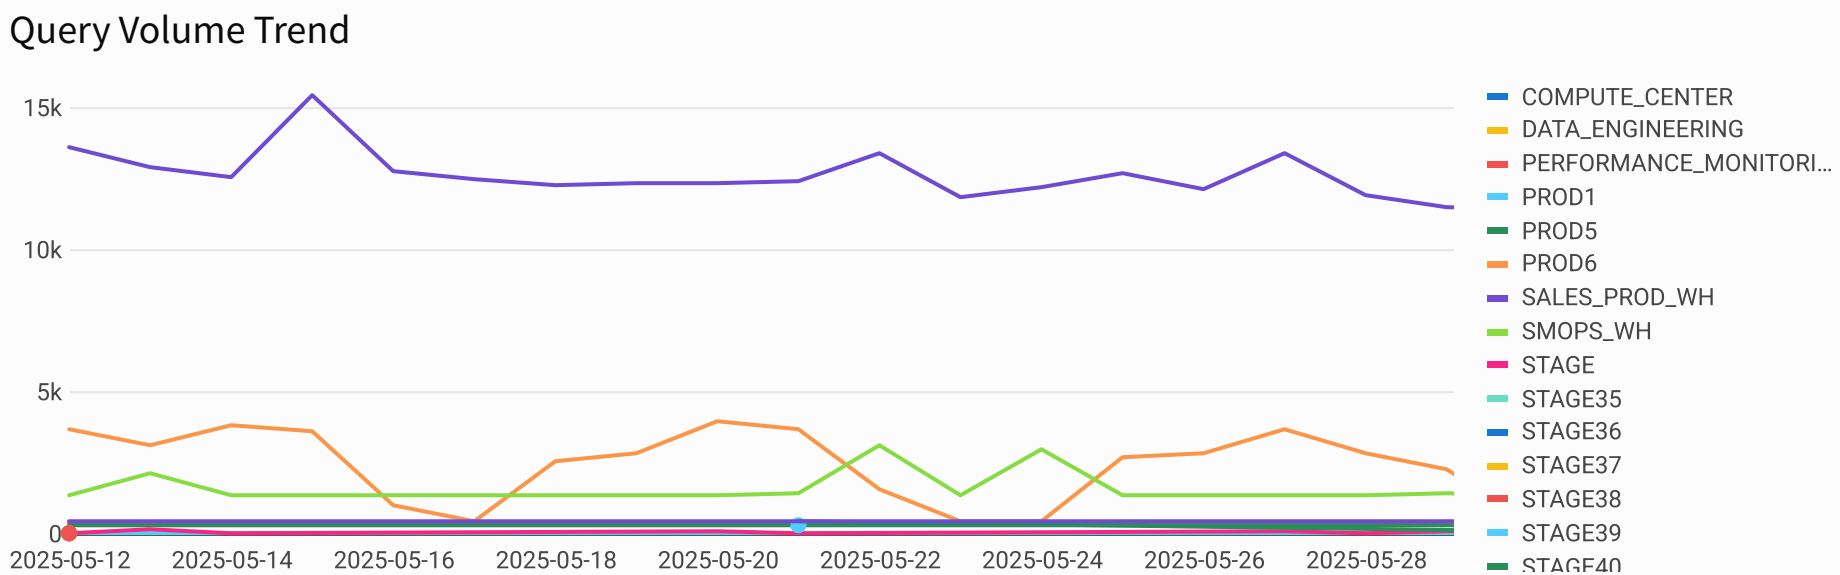


**Key Takeaways:**

**Spikes/dips:** Sharp peaks could correspond to a release or batch job; sudden drops might indicate an outage or changed user behavior.

**Predictability:** A regular weekly pattern (e.g., low on weekends, high on Mondays) tells you when to schedule maintenance or autoscaling policies.

---

## Average Execution Time for a Warehouse

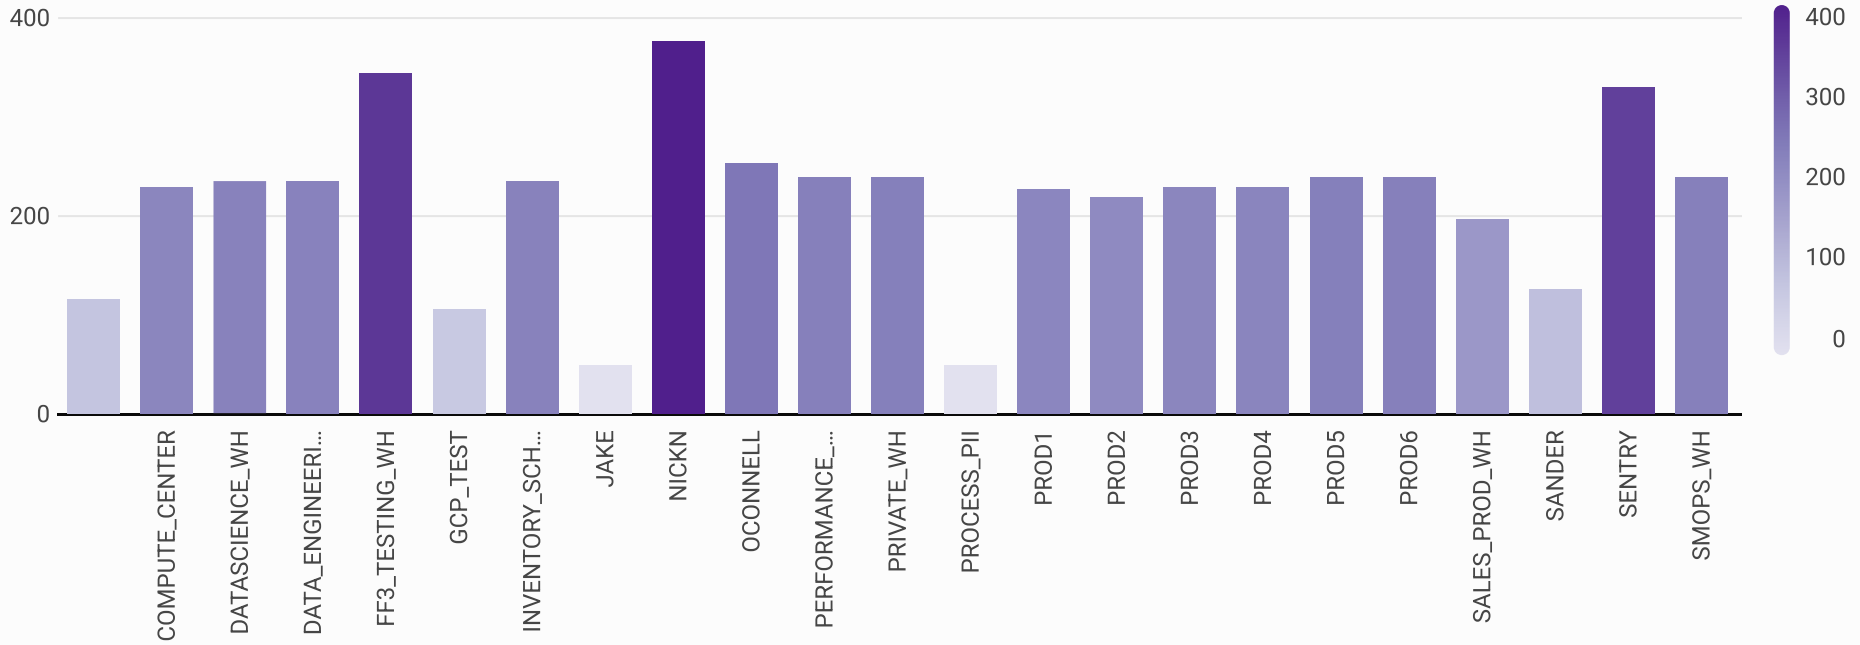


**Key Takeaways:**

**Performance bottlenecks:** Warehouses with high average times are slower—maybe under-powered, mis-configured, or querying very large datasets without proper clustering.

**Correlation with volume:** If a warehouse has both high volume and high latency, it's a prime candidate for optimization.

Data | Query History Analysis | Storage Insights | Events Insights | Credits Insights | Warehouse Efficiency Summary

Analysis Data

# Database Storage Growth and Insights

## Storage Usage Trends

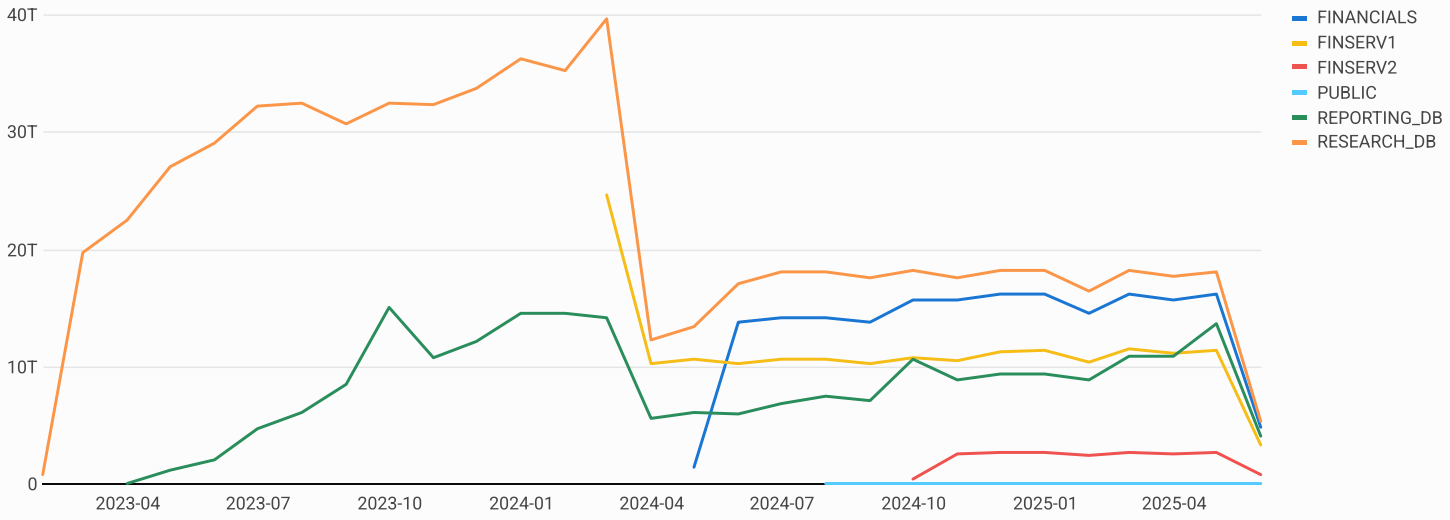


**Key Takeaways:**

**Rapid Growth:** Databases with steep slopes may need archiving or partitioning.

**Stability vs. Spikes:** Sudden jumps could indicate massive data loads or retention policy issues.

## Current Storage by Database

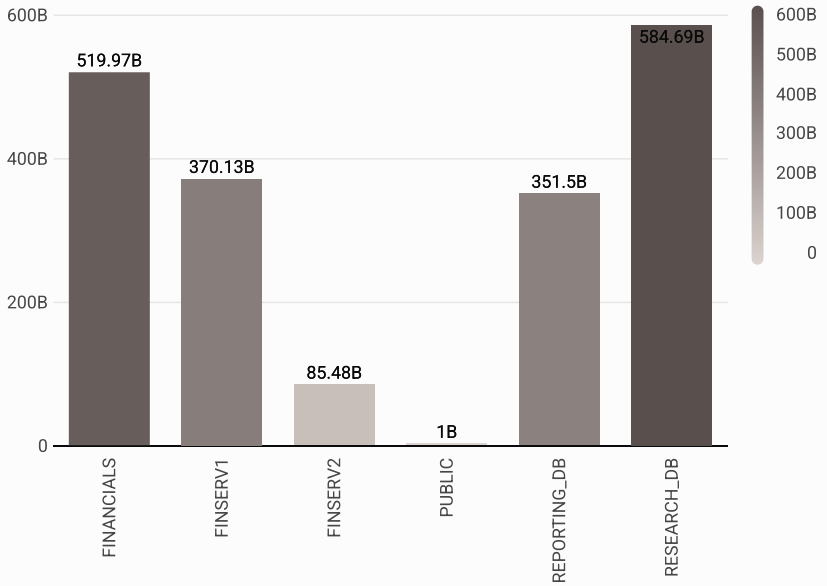


## Part of Total Storage for Each Database

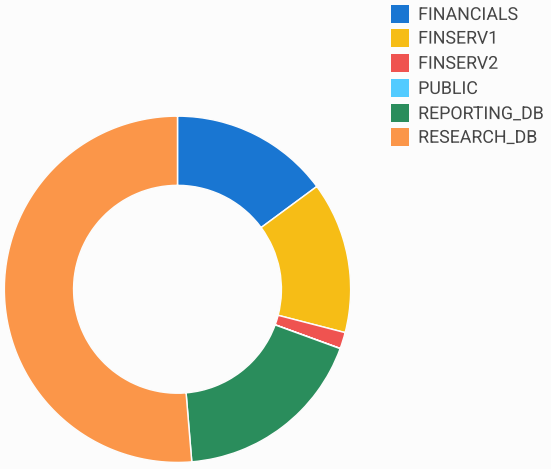


**Key Takeaways:**

**Cost Hotspots:** The top databases drive the majority of storage costs—candidates for compression or data cleanup.

**Resource Allocation:** Helps prioritize which data stores to optimize first.

Analysis Data

# Warehouse Event Patterns Dashboard

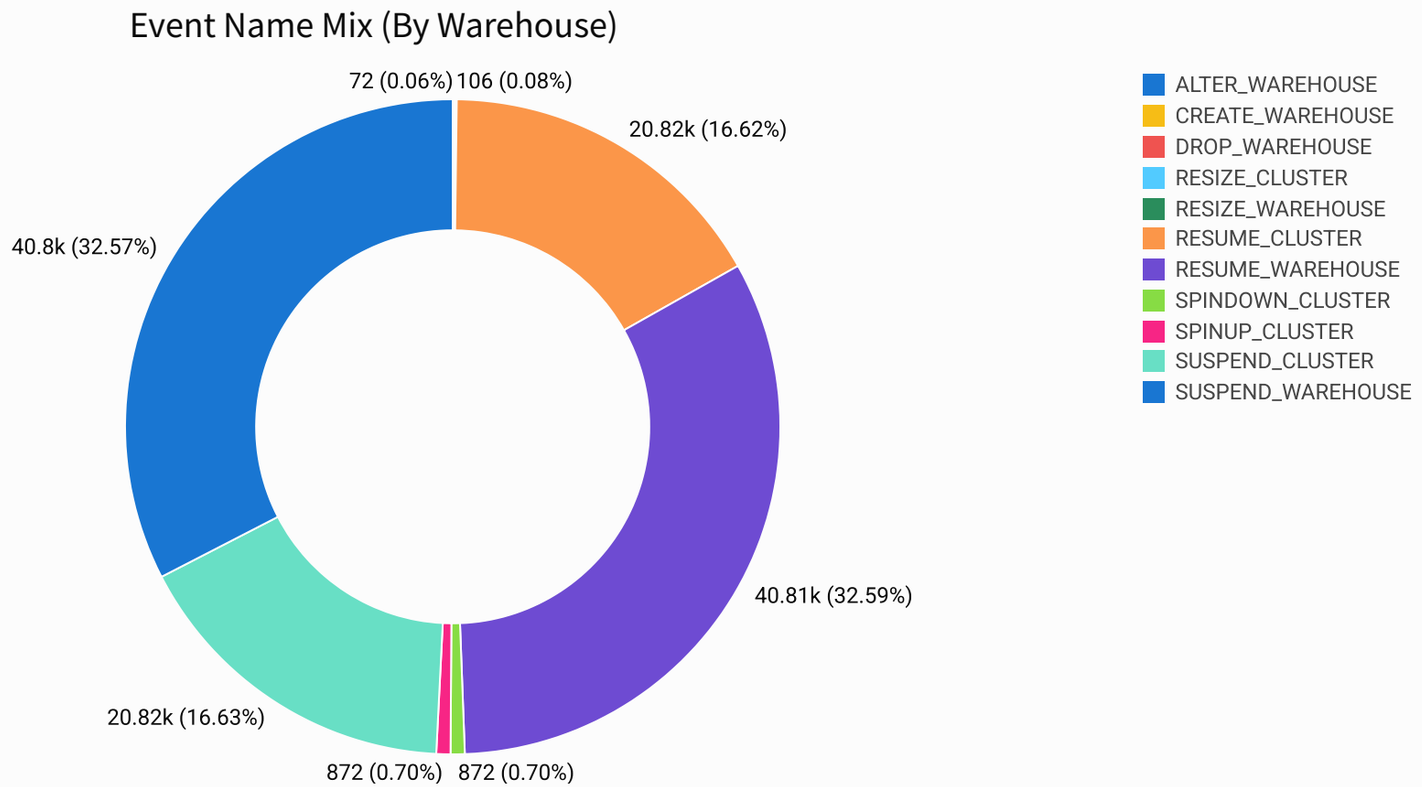


**Key Takeaways:**

We can see that 32.5% events are SUSPEND_WAREHOUSE and 32.5% are RESUME_WAREHOUSE, which indicate frequent autosuspend / autoresume cycles.

**Event Name**

(11) RESUME_WAREHOUSE, SUSPEND_WAREHOUSE, SUSPEND_CLUSTER, RESUME_CLUSTER, SPINDOWN_CLUSTER, SPINUP_CLUSTER, RESIZE_WAREHOUSE, ALTER_WAREHOUSE, RESIZE_C… ×

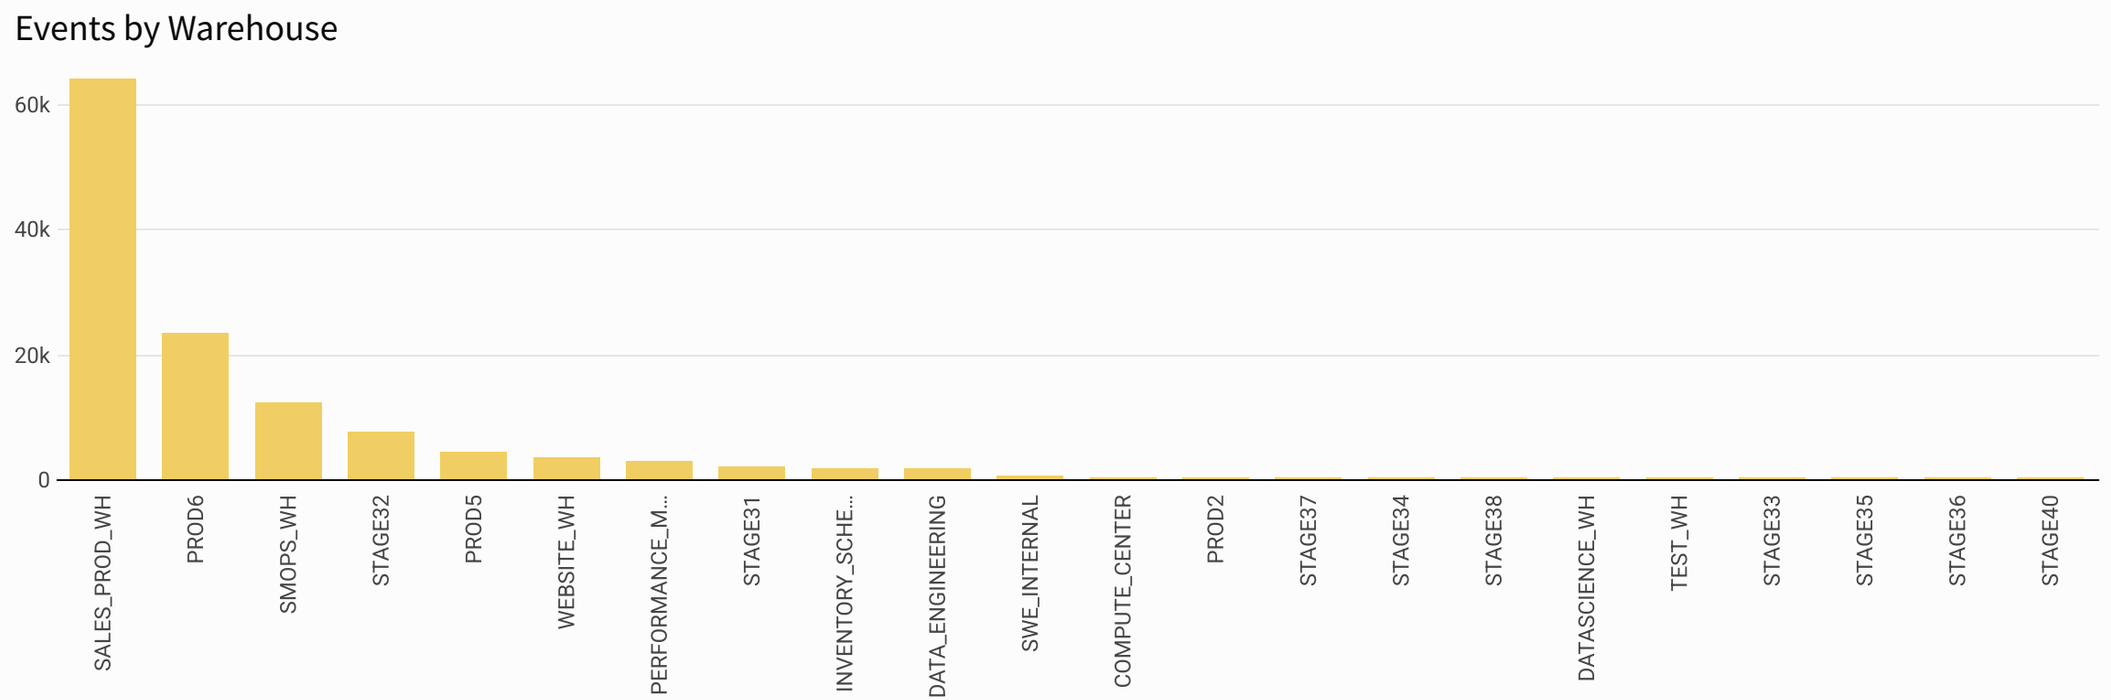


**Key Takeaways:**

Warehouse SALES_PROD_WH has 64,083 events. Its high event rate may signal oversizing or misconfigured autoscaling. We can investigate its usage patterns and right-size it.

## Recent Events

| Timestamp | Warehouse Name | Event Name | Event Reason |
|---|---|---|---|
| 2025-06-10 06:19:50 | SALES_PROD_WH | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND |
| 2025-06-10 06:19:50 | SALES_PROD_WH | SUSPEND_CLUSTER | WAREHOUSE_AUTOSUSPEND |
| 2025-06-10 06:19:50 | SALES_PROD_WH | SUSPEND_WAREHOUSE | WAREHOUSE_AUTOSUSPEND |
| 2025-06-10 06:11:41 | SALES_PROD_WH | RESUME_CLUSTER | WAREHOUSE_AUTORESUME |
| 2025-06-10 06:11:41 | SALES_PROD_WH | RESUME_WAREHOUSE | WAREHOUSE_AUTORESUME |

SUMMARY 125,247 rows – 5 columns

**Key Takeaways:**

Most recent events could be used to debug issues or figure out any automation misfires.

Analysis Data

# Credits Consumption & Cost-Efficiency Dashboard

## KPI: Sum of Credits Used

**5,458.46**

↓ 71.5% Jun vs May, 2025


## Sum of Credits Used Compute

**5,449.25**

↓ 71.5% Jun vs May, 2025

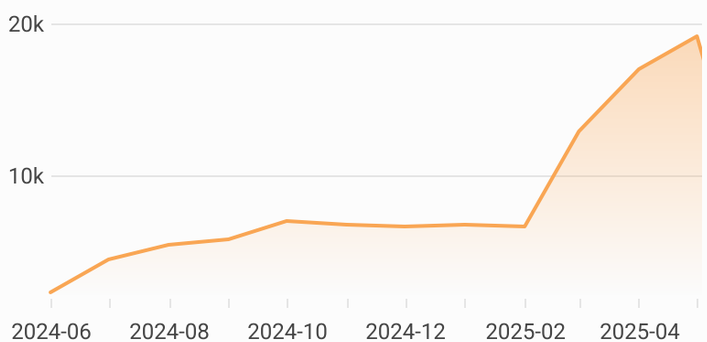


## Sum of Credits Used Cloud Services

**9.22**

↓ 71.5% Jun vs May, 2025

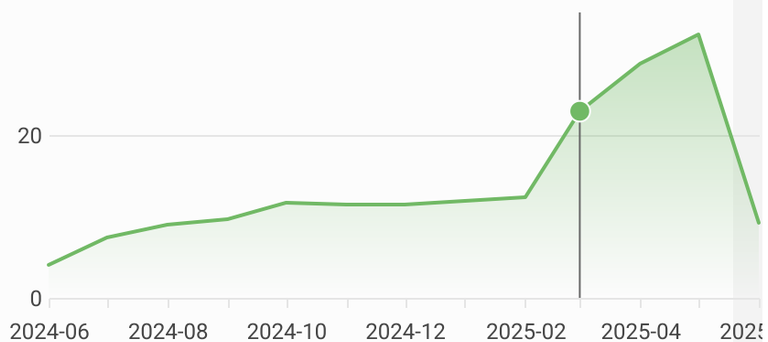


**Key Takeaways:** Shows the amount of credits we spent in the last 30 days, and a % change vs. the prior period. It can be used to indicate the value of our recent sizing adjustments or other implemented changes.

Also shows credits used for compute and cloud services separately to create showback or chargeback models to individual teams, enforcing financial accountability. (Used by FinOps teams)

**Start Time**

Last 30 days including today (05/12/2025 – 06/10/2025)

## Credits Used by Warehouse (Last 30 Days)

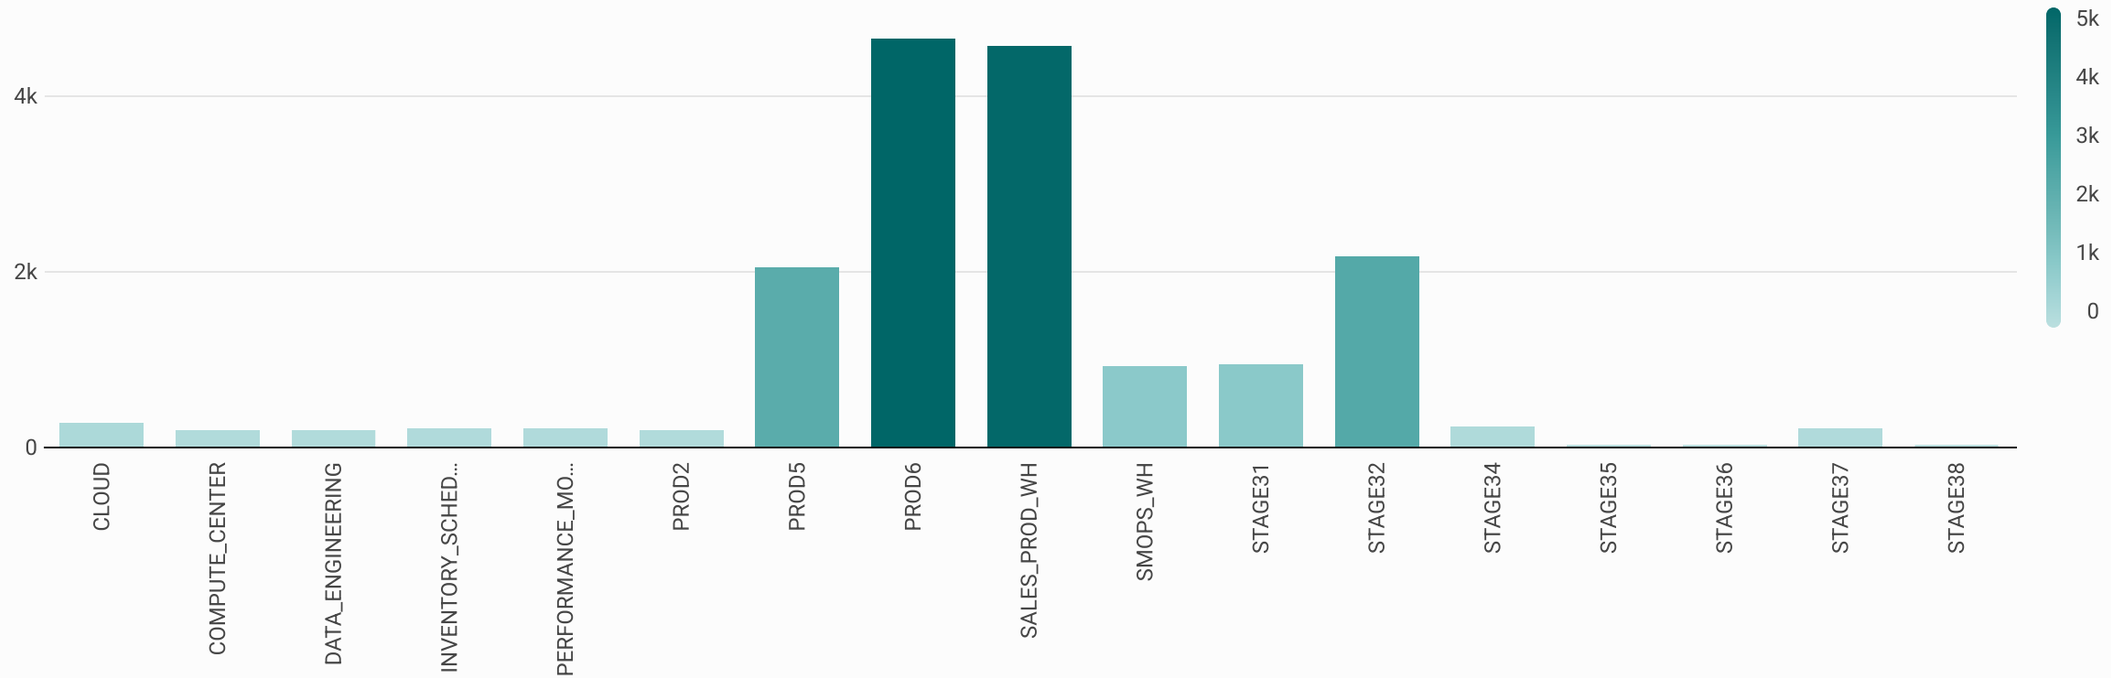


**Key Takeaways:**

**PROD6** and **SALES_PROD_WH** consumed **4491$** and **4408$** respectively, which means they are our top cost drivers. We should evaluate their auto-suspend settings to reduce idle spend.

## Credits over Time

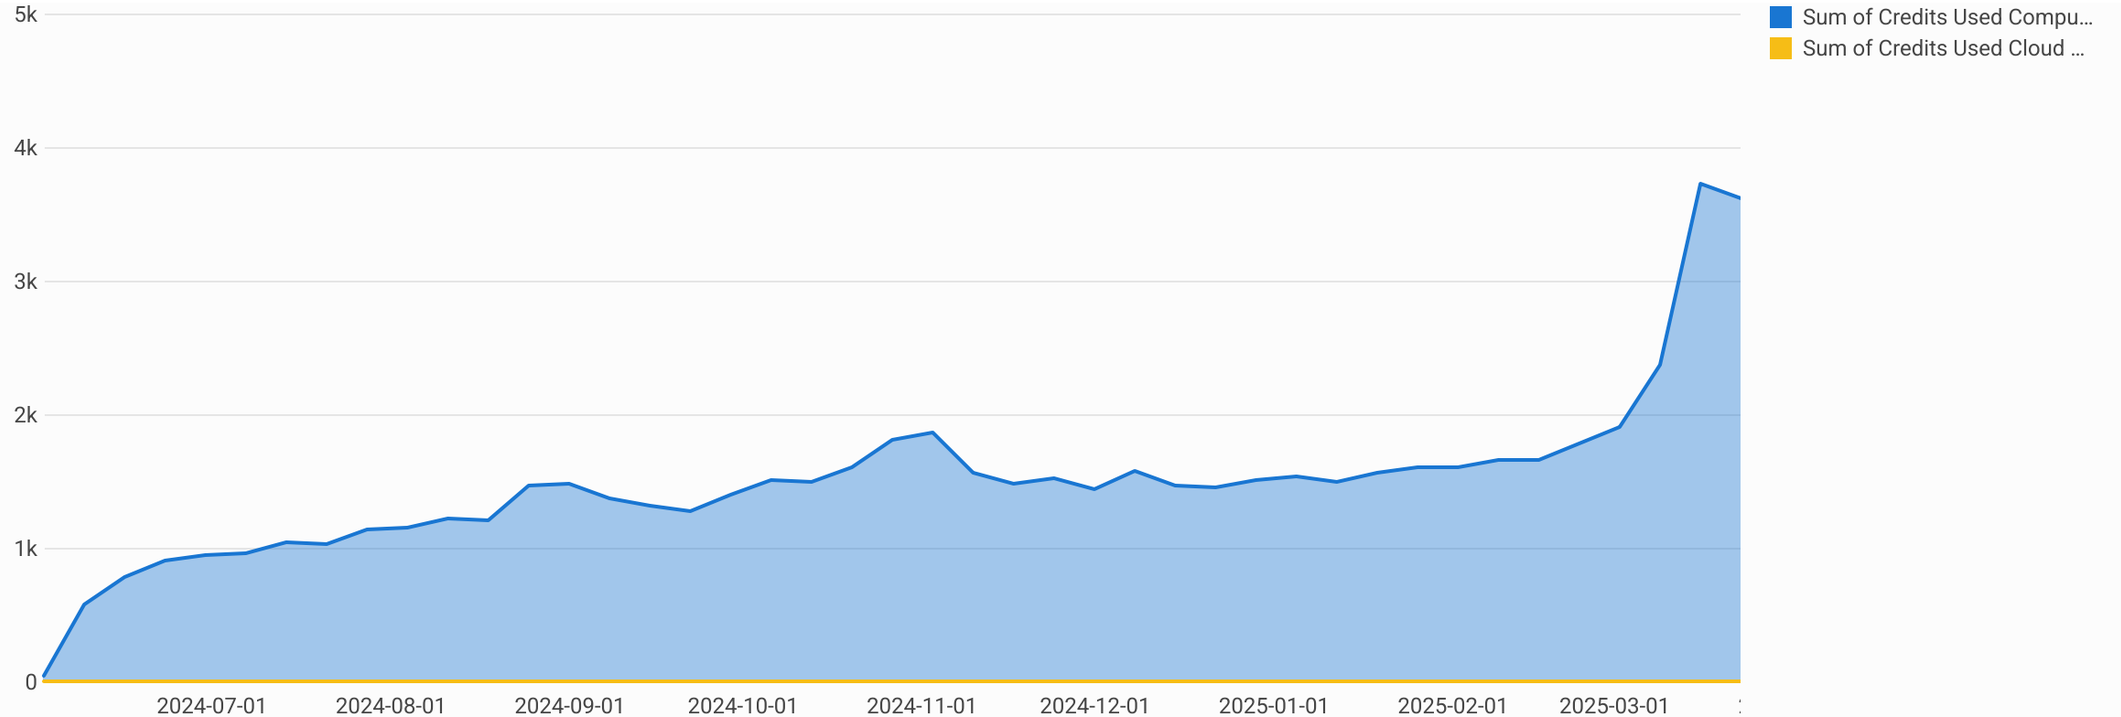


**Key Takeaways:**

**Peak usage weeks:** A week that shows a 20% spike in compute spend—might correlate to batch jobs.

**Shift in mix:** Cloud-service credits rose from 15% to 25% of total spend, suggesting increased use of external functions or data transfers.

**Downward trend:** After a mid-period peak, overall credits trended down 10%, indicating recent cost-optimization measures are working.

## Sum of Credits Used and Sum of Credits Used Compute by Warehouse Name

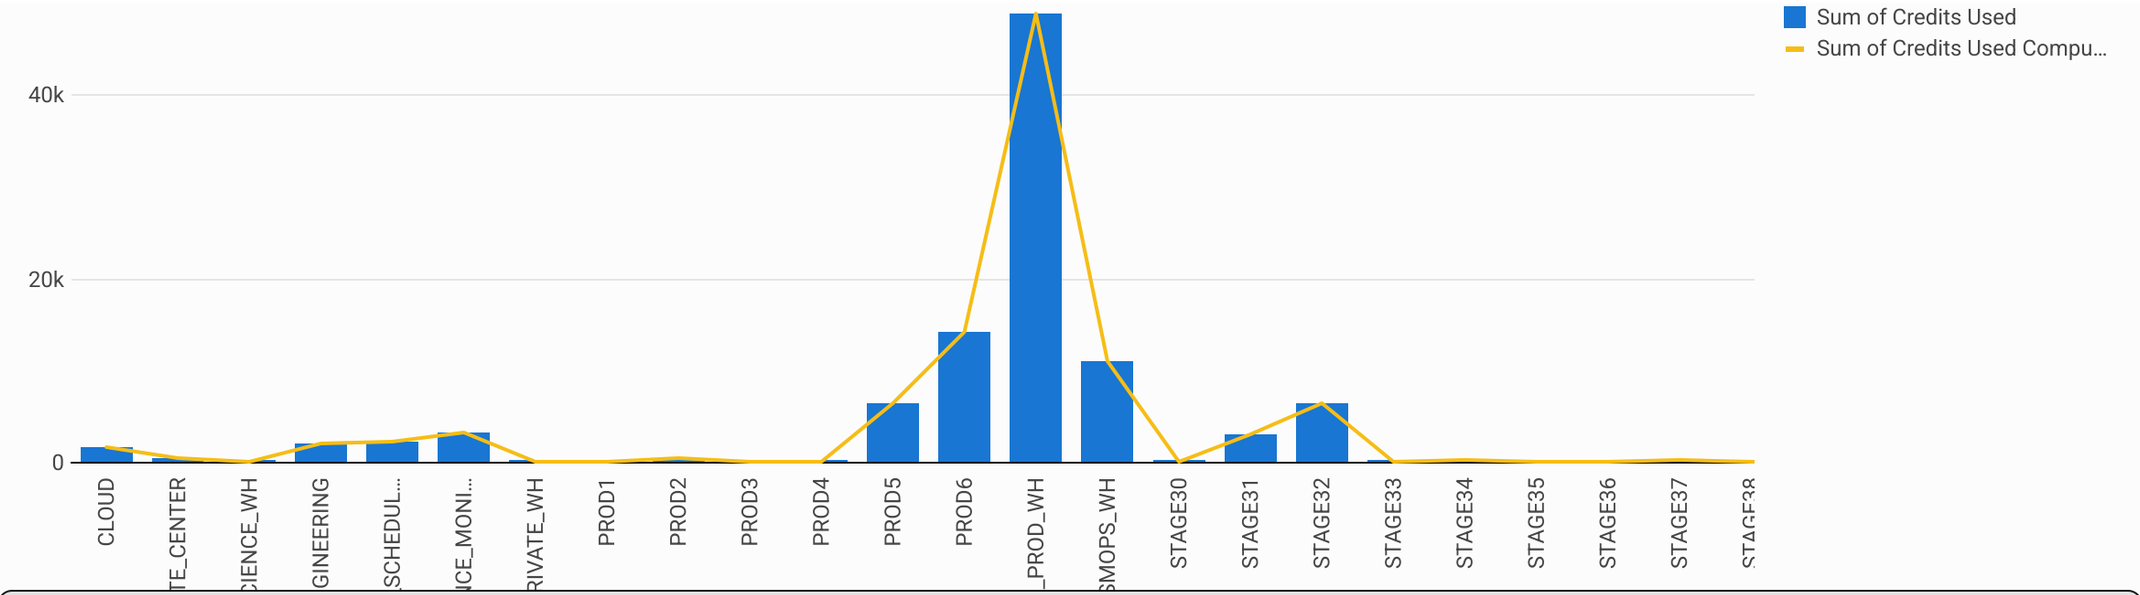


**Key Takeaways:**

**Heavy compute users:** Warehouses that consume almost all of total credits used for compute suggests CPU-intensive workloads.

**Balanced candidates:** Warehouses which show roughly 1:1 ratios in compute and cloud credits are good benchmarks for optimal sizing.

Data | Query History Analysis | Storage Insights | Events Insights | **Credits Insights** | Warehouse Efficiency Summary

Analysis Data

# Warehouse Efficiency & Cost-Performance Dashboard

## Queries per Compute Credit

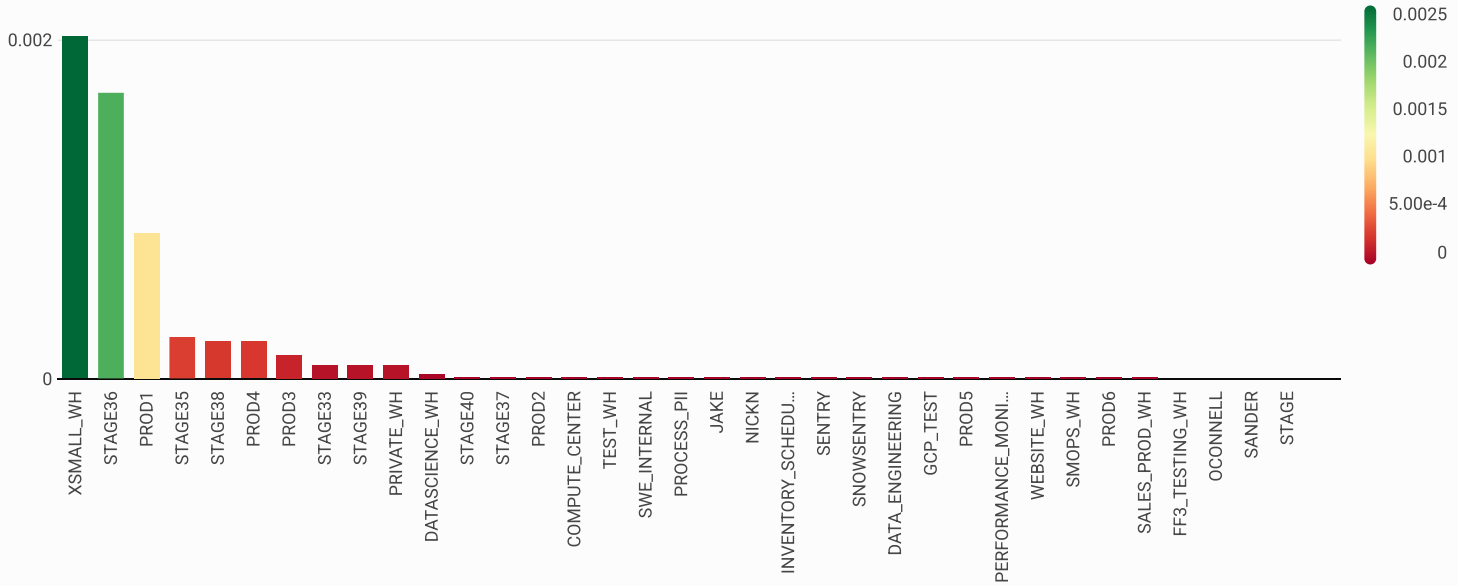


**Key Takeaways:**

Warehouses at the top deliver the most queries per compute credit, these are our cost-efficiency benchmarks. Those at the bottom likely need resizing or query optimization.

## Cost per Query

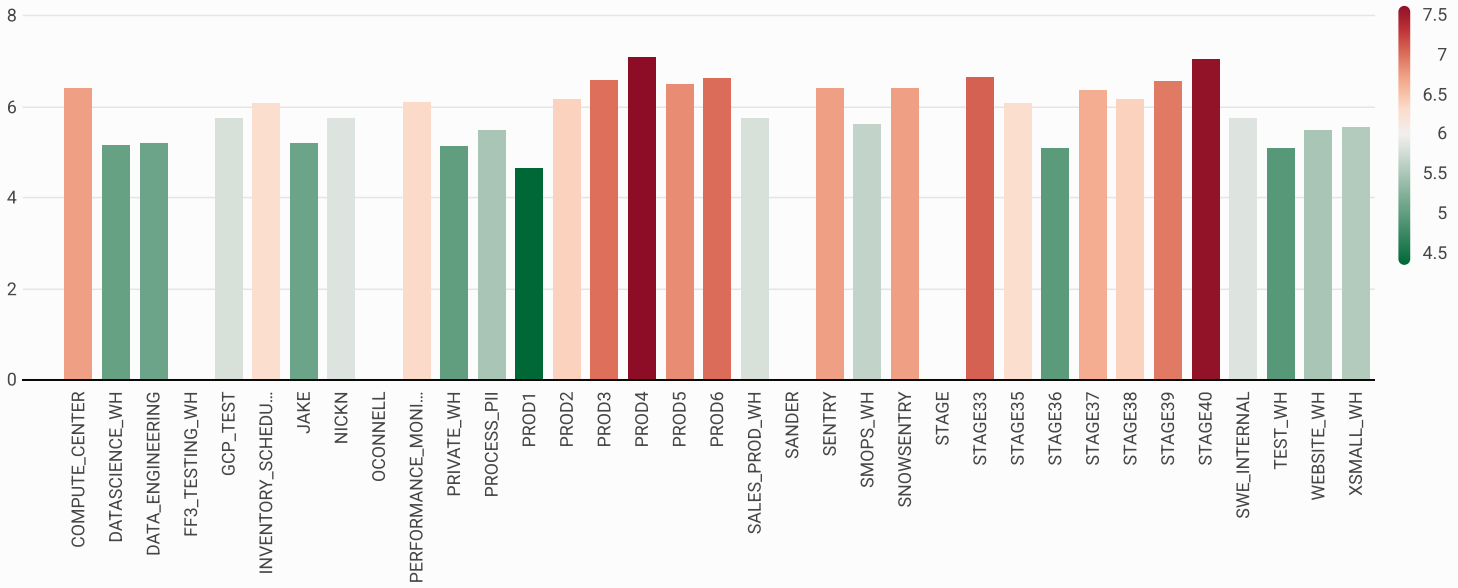


**Key Takeaways:**

**High-cost outliers:** Warehouses at the top are spending significantly more credits per query and these can becandidates for query optimization or down-sizing.

**Low-cost benchmarks:** Warehouses at the bottom deliver queries at a much lower cost, which are good reference configurations to replicate.